Sardar Balbeer Singh

Principal Data Architect | Data Platforms & Engineering

Hyderabad, India | +91 89777 47847 | balbeer.bedi9@gmail.com | linkedin.com/in/balbeer-singh-bedi

PROFILE

Data architect with 16 years of experience building data platforms for banking, insurance, SaaS, mortgage, and membership organizations. My work has ranged from hands-on data engineering to architecture and delivery leadership. Recent projects include an AWS real-estate data platform, a Microsoft Fabric analytics lakehouse, and a large-scale SaaS data migration.

TECHNICAL SKILLS

Cloud: AWS, Microsoft Azure, Google Cloud Platform

Data platforms: Microsoft Fabric, Databricks, Snowflake, BigQuery, Synapse, Teradata

Engineering: Python, SQL, PySpark, Apache Spark, Airflow, Kafka, Dataflow, Hive, PostgreSQL, MongoDB

Architecture: Lakehouse and data warehouse design, ETL/ELT, CDC, data migration, data quality, CI/CD

PROFESSIONAL EXPERIENCE

Principal Data Architect | V2Solutions

March 2025–Present | Mumbai, India

CrossCountry Mortgage | Real-estate data platform

- Built incremental AWS pipelines for millions of property records from ATTOM and CoreLogic. Watermark-based processing replaced full-table scans and reduced compute costs by 80%.
- Developed a Spark-based 15-day rule to distinguish market events from vendor corrections.
- Designed nested and flat MongoDB models for property history and fast search, with address standardization across vendor feeds.

AVIXA | Microsoft Fabric analytics

- Brought Salesforce, Sage Intacct, and Pardot data into a Microsoft Fabric lakehouse, using a medallion architecture to organize ingestion, transformation, and reporting.
- Delivered Power BI reporting for customer value, sales, and audience engagement. Reconciled CRM pipeline with recognized ERP revenue within a 2% tolerance.

Thrive TRM | SaaS migration validation

- Built Python and PostgreSQL checks to compare legacy and target data across 134+ tables for 270+ customer migrations, covering 1.2M+ people, 193K+ companies, and 99K+ candidacy relationships.
- Added field-level discrepancy reports and schema-drift checks to find issues before cutover; the validation work supported zero-defect go-lives.

Architect, Data Solutions | Apexon

July 2022–March 2025

- At Crux Informatics, built configurable Python and YAML data workflows orchestrated with Airflow, added reusable production QA checks, and supported Databricks and Snowflake prototypes.
- Prototyped event-triggered GCS workflows with schema evolution and Type 1/Type 2 Delta tables.
- At One Dealer Lane, designed Azure Synapse pipelines from Azure SQL through ADLS into Delta tables, along with data models and a Git-based deployment approach.

Cloud Architect / Data Engineer | HSBC Technology India

August 2016–July 2022 | Hyderabad, India

- Helped move global operations data to GCP, including near-real-time CDC and Kafka feeds for 20 EOPS tables into GCS and BigQuery, plus historical loads from DB2.
- Built Hadoop and Hive data-lake reporting for revenue-leakage analysis, and later contributed to a GCP warehouse migration and forecasting work with data scientists.

Senior System Engineer | IBM India

August 2010–August 2016 | Client: Nationwide Insurance

- Built and tested ETL mappings for Commercial Velocity, Smart Ride, and TTS, using Teradata utilities to load and extract data for enterprise reporting.
- Worked on a 100+ TB Teradata warehouse integrating data from more than 48 source systems.

EDUCATION

Bachelor of Engineering, Electronics & Communication Engineering

GNDEC, Bidar | Visvesvaraya Technological University | 2007

CERTIFICATIONS

- Google Cloud Professional Cloud Architect
- Professional Scrum Master I (Scrum.org)
- Teradata 12 Basics
- Big Data University